

COMN - Computer Communications and Networks

Assignment 2: Implement & Analyse Sliding Window Protocols

Michio Honda

With acknowledge to Andrew Ferguson and Jon Larrea

School of Informatics

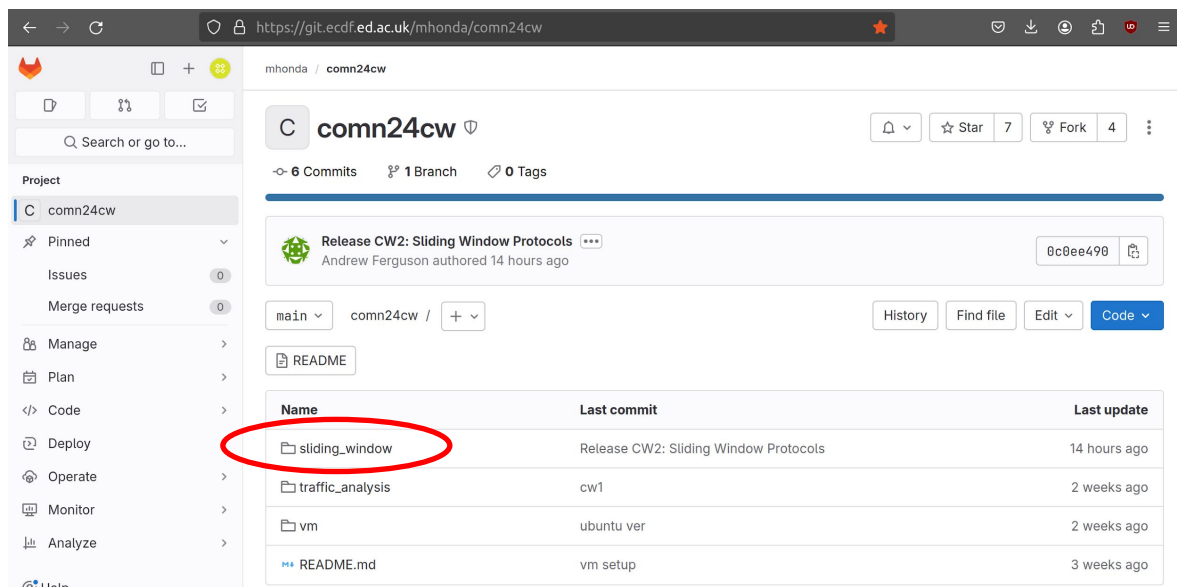
University of Edinburgh

14/02/2025



Assignment 2 Released

- Due Monday 31th March, midday (1200 UTC / GMT)
- Available at <https://git.ecdf.ed.ac.uk/mhonda/comn25cw>
 - Linked from the Piazza announcement will be emailed



mhonda / comn24cw

comn24cw

6 Commits 1 Branch 0 Tags

Release CW2: Sliding Window Protocols
Andrew Ferguson authored 14 hours ago

8c8ee490

main comn24cw +

History Find file Edit Code

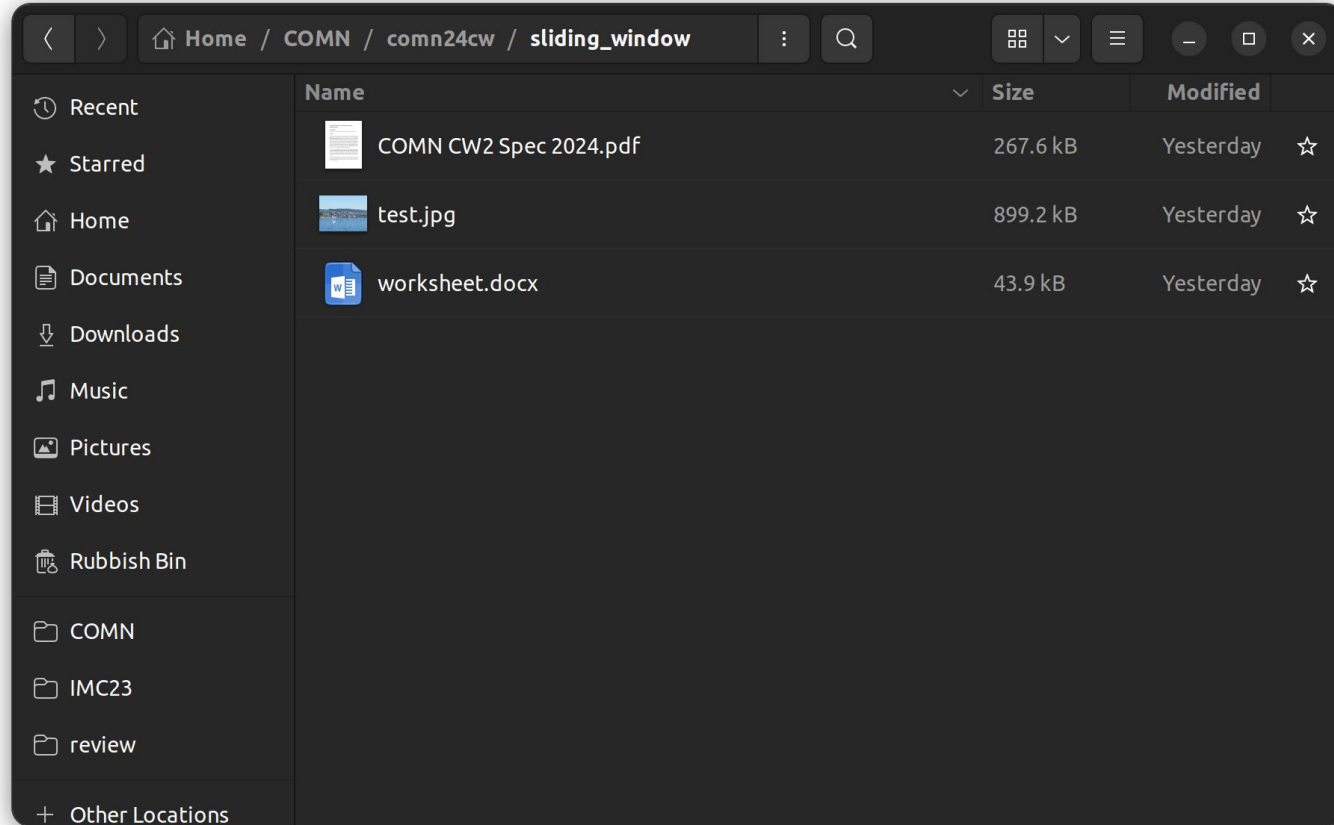
README

Name	Last commit	Last update
sliding_window	Release CW2: Sliding Window Protocols	14 hours ago
traffic_analysis	cw1	2 weeks ago
vm	ubuntu ver	2 weeks ago
README.md	vm setup	3 weeks ago

Assignment 2 Released

- "I have already cloned the git repo (e.g. to setup the VM, or attempt CW1)"
 - `cd comn25cw`
 - `git pull`
 - `<enter your username / password>`
- "I have not yet cloned the git repo"
 - `git clone https://git.ecdf.ed.ac.uk/mhonda/comn25cw.git`
 - `<follow the VM setup instructions in comn24cw/vm>`

Assignment 2 Released



Coursework Overview

- Goal
 - Implementation and evaluation of three end-to-end reliable data transfer protocols
- Assessment: 30% of course mark
 - Part 1 (4.2%): rdt1.0
 - Part 2 (8.6%): rdt 3.0 (Stop-and-Wait)
 - Part 3 (8.6%): Go-back-N
 - Part 4 (8.6%): Selective Repeat + iPerf experiment

Reliable Data Transport Protocols

- Why do we need a reliable data transport protocol?
- Consider TCP (reliable) vs UDP (unreliable)

TCP




UDP



Reliable Data Transport Protocols

- TCP: Ensures each packet is received by the receiver
- UDP: sends data without caring what the receiver is doing (or even if there is one at all)
- For this assignment: "reliable" == no **lost** packets
 - Data corruption is a separate concern
 - We assume no corruption can occur
 - In the real world checksums can detect and correct errors

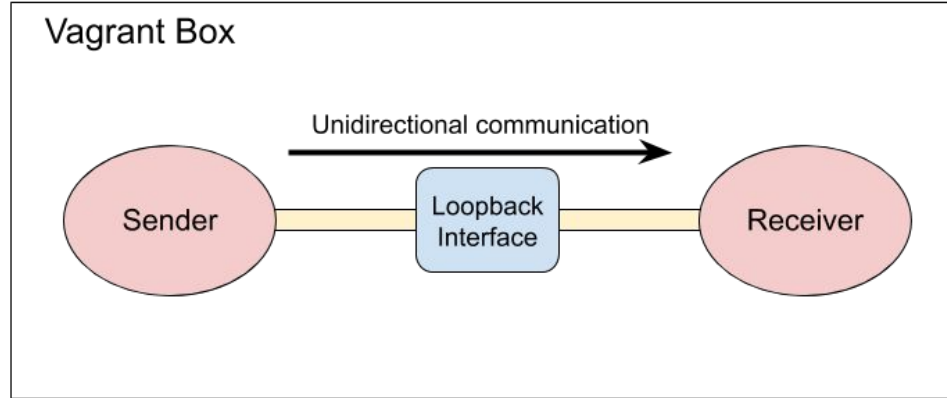
Conceptual Structure



■ Linux Traffic Control (The tc utility)

- Allows you to modify the packet scheduler for a given interface
- Configuration of interface characteristics (bandwidth, delay, loss)
- Command-line program: tc

Conceptual Structure



- Sender
 - Reads the file and breaks it into a number of packets
 - Sends the packets to a receiver over the loopback interface which has the forwarding rules modified
- Receiver
 - Receives the packets; extracts data in the packets; and saves the data in a file

Header format

- The following formats should be used across all parts

- Exception: no ACK packets needed in part 1

- Data packet

- (Sender to Receiver)

0	1	2	3 ~ up to 1026
16-bit sequence number		8-bit EoF flag	Data

- ACK packet

- (Receiver to Sender)

0	1
16-bit sequence number	

Useful Links

- These two links are for simulators for the go-back N / selective repeat protocols that you need to implement.
- They are one of the best ways to improve your understanding of how the algorithm works and how it handles edge cases.
- https://media.pearsoncmg.com/aw/ecs_kurose_compnetwork_7/cw/content/interactiveanimations/go-back-n-protocol/index.html
- https://media.pearsoncmg.com/aw/ecs_kurose_compnetwork_7/cw/content/interactiveanimations/selective-repeat-protocol/index.html

Useful tools: iPerf

- **iPerf** is a tool used to measure network performance measurement in terms of throughput and latency.

Client

```
openair@openair-1:~$ iperf -c 192.168.4.5 -i1 -t10
-----
Client connecting to 192.168.4.5, TCP port 5001
TCP window size: 85.0 KByte (default)
-----
[ 3] local 192.168.4.10 port 34562 connected with 192.168.4.5 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 1.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 1.0- 2.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 2.0- 3.0 sec   11.1 MBytes 93.3 Mbits/sec
[ 3] 3.0- 4.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 4.0- 5.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 5.0- 6.0 sec   11.1 MBytes 93.3 Mbits/sec
[ 3] 6.0- 7.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 7.0- 8.0 sec   11.1 MBytes 93.3 Mbits/sec
[ 3] 8.0- 9.0 sec   11.2 MBytes 94.4 Mbits/sec
[ 3] 9.0-10.0 sec   11.1 MBytes 93.3 Mbits/sec
[ 3] 0.0-10.0 sec   112 MBytes 93.9 Mbits/sec
```

Server

```
openair@openair-1:~$ iperf -s
-----
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
-----
```

Useful tools: iPerf

```
iperf -c 192.168.4.5 -i1 -t10
```

```
iperf -c 192.168.4.5 -i1 -n 30MB
```

```
iperf -c 192.168.4.5 -i1 -F test.jpg -M 1KB
```

- -c → Receiver IP address
- -i → Interval, seconds between periodic bandwidth reports
- -t → time in seconds to transmit for (default 10 secs)
- -n → number of bytes to transmit (instead of -t)
- -F → input the data to be transmitted from a file
- -M → set TCP maximum segment size

Useful tools: tcpdump

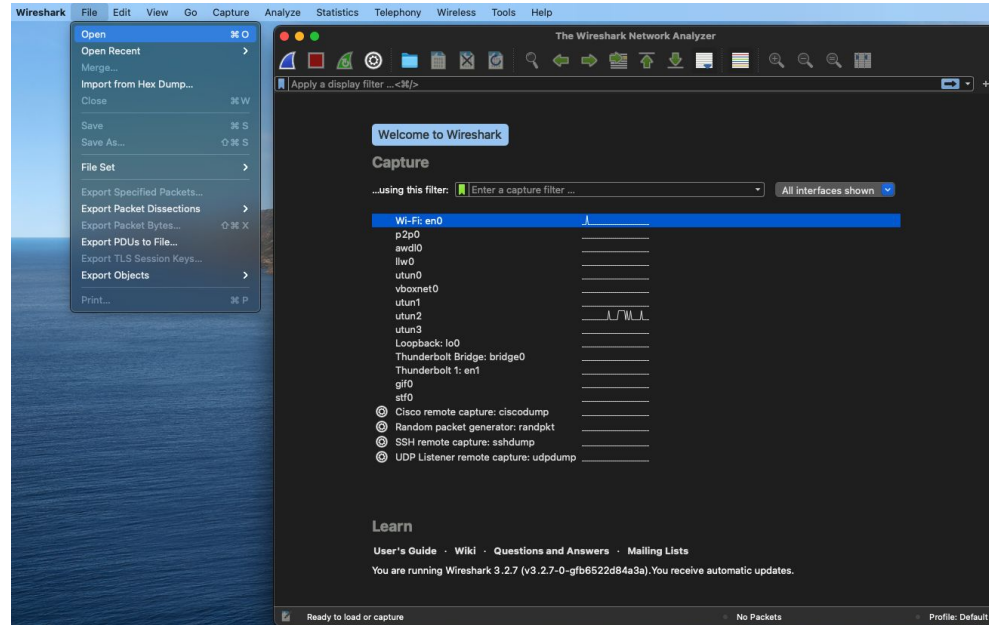
- **tcpdump** is a data-network packet analyser tool. (Command line version of **Wireshark**).
- Due to the lack of a GUI in the COMN VM, we cannot run **Wireshark** in the VM. We can however obtain the packet capture trace from inside the VM using **tcpdump**, then examine it outside the VM on our host machine using **Wireshark**.

```
sudo tcpdump -i lo -w out.pcap
```

- `-i` → Interface (loopback).
- `-w` → Write the raw packets to a file.

Useful tools: Wireshark

- **Wireshark** is an open-source packet analyser tool that is used to capture network packets to understand and troubleshoot network behaviour.



Useful tools: Wireshark

The screenshot shows the Wireshark interface with a packet capture list and packet details pane. Annotations highlight key fields for analysis:

- 1. Packet Number:** Points to the 'No.' column in the packet list.
- 2. Timestamp:** Points to the 'Time' column in the packet list.
- 3. Source IP Address:** Points to the 'Source' column in the packet list.
- 4. Destination IP Address:** Points to the 'Destination' column in the packet list.
- 5. Protocol:** Points to the 'Protocol' column in the packet list.
- 6. Source and Destination Ports:** Points to the 'Info' column in the packet list.

The packet list shows several packets, with packet 664 selected. The packet details pane shows the following information:

- Frame 664 (68 bytes on wire, 68 bytes captured)
- Linux cooked capture
- Internet Protocol, Src: 10.10.0.133 (10.10.0.133), Dst: 10.10.0.107
- Transmission Control Protocol, Src Port: ssh (22), Dst Port: 33048, Seq=15325, Ack=4949, Win=2140, Len=0, TSV=369138, TSEQ=2804771523

The packet bytes pane shows the raw data in hexadecimal and ASCII format.

Useful tools: Wireshark

The screenshot shows the Wireshark interface with a packet capture of an SSH session. The packet list pane shows a list of packets, and the packet details pane shows the details of the selected packet (Frame 664). The packet details pane is expanded to show the 'Transmission Control Protocol' section, which includes fields for 'Sequence number: 16001' and 'Window size: 2140'. Two callout boxes are present: one pointing to the 'Sequence number' field labeled '7. TCP sequence number' and another pointing to the 'Window size' field labeled '8. TCP window size'.

Frame 664 (68 bytes on wire, 68 bytes captured)

Linux cooked capture

Internet Protocol, Src: 10.10.0.133 (10.10.0.133), Dst: 10.10.0.107 (10.10.0.107)

Transmission Control Protocol, Src Port: ssh (22), Dst Port: ssh (22), Seq: 16001, Ack: 5069, Len: 0

Source port: ssh (22)

Destination port: 33048 (33048)

[Stream index: 1]

Sequence number: 16001 (relative sequence number)

Acknowledgement number: 5069 (relative ack number)

Header length: 32 bytes

Flags: 0x10 (ACK)

Window size: 2140

Checksum: 0x49dd [Verify]

Options: (12 bytes)

[SEQ/ACK analysis]

7. TCP sequence number

8. TCP window size

Miscellaneous

- Some useful Python libraries for the assignment:
 - `sys`
 - `socket`
 - `math`
 - `time`
 - `Thread` from `threading`
 - `Lock` from `threading`

Miscellaneous

- Some useful Python libraries for the assignment:

- sys
- socket
- math
- time
- Thread from threading
- Lock from threading

All of these are installed by default (in Python standard library).
Do not install your own modules – you will FAIL.

Design choices for Part 3 and 4

- Both sender and receiver are implementable without multithreading
 - Definitely no need for multithreading at the receiver side
 - Multithreading may be useful for sender implementation
 - e.g. threads for (a) sending, (b) receiving (ACKs), (c) checking timeouts
- Use non-blocking socket for non-multithreaded implementation
 - `setblocking(0)` and `select()`
 - Why? Whilst waiting to receive an ACK (or a timeout) you may need to send another packet
- **Many** design choices for the sender are possible

How is CW2 marked?

1. We (the markers) manually run your code.
2. We check that:
 - The file is transferred correctly (byte level).
 - Your code did not crash.
 - Your code did not hang.
3. We run **tcpdump** to capture the network traffic caused by your code.
4. We use **WireShark** to check the captured packets to see if you have implemented the protocols correctly.
5. The worksheet is marked more conventionally using a marks scheme.

Common Mistakes

1. Your code does not run or crashes.
 - **DO NOT INSTALL ANY ADDITIONAL PYTHON PACKAGES**
 - If in doubt, create a fresh version of the vagrant VM, and check there.
2. Your codes does not exit.
 - e.g. receiver waits for the final packet with EoF flag, that packet is lost
 - More marks lost for code that does not exit than incorrectly transferred file
3. Timeouts **are in milliseconds**
 - Not seconds!

Less Common Mistakes

1. Please don't try using TCP
 - (We will notice)

Getting Help

- Piazza
 - Ask general questions to the entire class
 - Ask code-specific questions to the instructors only
- Weekly Q+A Sessions
 - Will be focusing on CW2 for the next four weeks

Questions?

Optional: slides on multithreading (*borrowed* from Operating Systems lectures)

Appendix: Multithreading

What is multithreading?

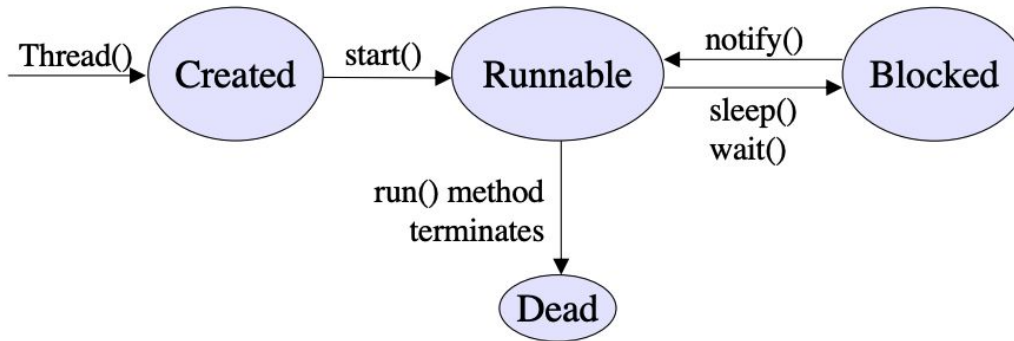
- Multithreading is similar to multi-processing
- A multi-processing OS can run several processes at the same time
 - Each process has its own address/memory space
 - Separate processes do not have access to each other's memory space
- In a multithreaded application, there are several points of execution **within the same memory space**
 - Each point of execution is called a thread
 - Threads share access to memory

Thread support in Python

- Python threading allows you to have different parts of your program running concurrently.
- Threads are represented by a Thread object
 - A thread object maintains the state of the thread
 - It provides control methods such as start, run, sleep, join
- When an application executes, the main method is executed by a single thread
 - If the application requires more threads, the application must create them

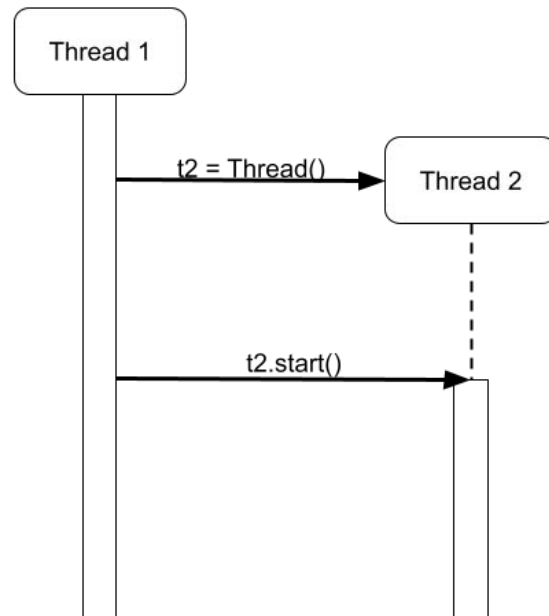
Threads States

- Threads can be in one of four states as shown in the figure below
- A thread's state changes based on:
 - Control methods such as start, sleep, wait, notify
 - Termination of the run method



How does a thread run?

- The thread class has a `run()` method
 - `run()` is executed when the thread's start method is invoked
- The thread terminates if the run method terminates
 - `run()` method often has an endless loop to prevent thread termination
- One thread starts another by calling its start method



Creating your own thread

- The Python standard library provides **threading**, which contains all the primitives you'll need for this assignment.

```
import threading
import time

def thread_function(name):
    print('Hello from ' + name)

    for i in range(5):
        print(i)
        time.sleep(1)

if __name__ == '__main__':
    print('Before creating the thread')
    your_thread = threading.Thread(target=thread_function, args=('My Thread', ))
    print('Before running thread')
    your_thread.start()

    print('Wait for the thread to finish')
    your_thread.join()
    print('End')
```


```
jon@macbook:~/Desktop$ python test.py
Before creating the thread
Before running thread
Hello from My Thread
Wait for the thread to finish
0
1
2
3
4
End
```

Synchronization: Critical Sections/Mutual Exclusion

- Sequence of instructions that may get incorrect results if executed simultaneously are called **critical sections**
- We also use the term **race condition** to refer a situation in which the results depends on timing
- **Mutual exclusion** means “not simultaneous”
 - $A < B$ or $B < A$
 - We don't care which
- Forcing mutual exclusions between two critical section executions is sufficient to ensure the correct execution – guarantees ordering
- One way to guarantee mutually exclusive executions is using **locks**

Critical sections

→ is the "happens-before" relation



When do critical sections arise?

- One common pattern:
 - read-modify-write of a shared value (variable) in code that can be executed concurrently
- Shared variable:
 - Globals and heap-allocated variables
 - NOT local variables (which are on the stack)

Example: Shared bank account

- Suppose we have to implement a function to withdraw money from a bank account:

```
int withdraw(account, amount) {  
    int balance = get_balance(account);    // read  
    balance -= amount;                      // modify  
    put_balance(account, balance);         // write  
    spit out cash;  
}
```

- Now suppose you and your partner share a bank account with a balance of \$100.00
 - What happened if you both go to separate ATM machines, and simultaneously withdraw \$10.00 from the account?

Example: Shared bank account

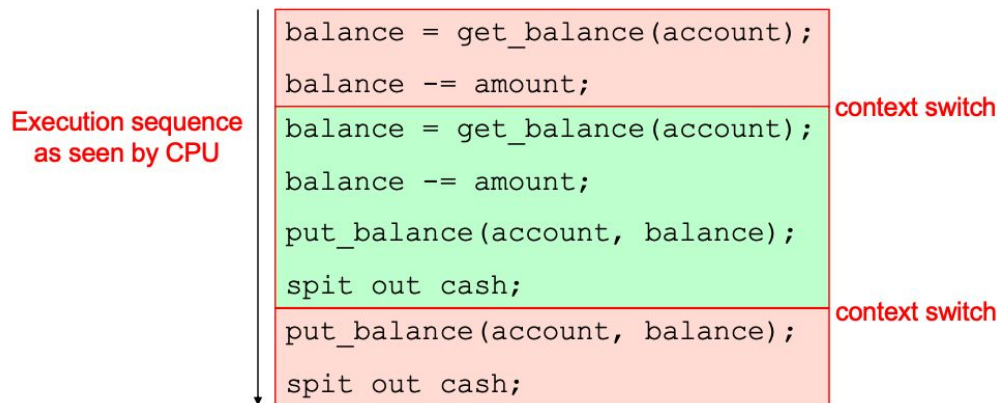
- Assume the bank's application is multi-threaded
- A random thread is assigned a transaction when that transaction is submitted

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    spit out cash;  
}
```

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    spit out cash;  
}
```

Interleaved schedules

- The problem is that the execution of the two threads can be interleaved, assuming preemptive scheduling:



- What's the account's balance after this sequence?
 - Who's happy, the bank or you?

Locks

- A lock is a memory object with two operations:
 - **acquire()**: obtain the right to enter the critical section
 - **release()**: give up the right to be in the critical section
- **acquire()**: prevents progress of the thread until the lock can be acquired

